

FastCite: A RAG-Based AI Study Assistant Chatbot with Hybrid Search and Intelligent Context Selection for Academic Learning

Aden Sial¹, Inam Ullah Shaikh¹, and Emaan Fatima¹

National University of Computing and Emerging Sciences, Islamabad, Pakistan
{22i-1313, 22i-0857, 22i-0869}@nu.edu.pk

Abstract. The proliferation of digital learning materials has created unprecedented opportunities for students to access vast collections of academic documents, yet finding relevant information and getting personalized study assistance remains challenging. This paper presents FastCite, an AI-powered study assistant chatbot that helps students learn from their uploaded course materials through intelligent document retrieval and context-aware question answering. Our system addresses key limitations of existing study tools by implementing a hybrid search mechanism that combines semantic understanding with keyword matching, enabling students to ask questions in natural language and receive accurate, well-sourced answers from their study materials. FastCite employs the allmpnet-base-v2 embedding model for semantic understanding, Qdrant vector database for efficient similarity search, and Google Gemini 2.5 Flash Lite for conversational answer generation. We introduce a weighted hybrid search approach (60% vector, 40% keyword) that significantly improves retrieval accuracy compared to pure semantic or keyword-based methods, allowing students to find information whether they use exact terminology or describe concepts in their own words. Additionally, we implement an LLM-based context selection mechanism that intelligently filters retrieved study materials to provide the most relevant information, ensuring students receive focused, helpful responses. Through comprehensive experiments and ablation studies, we demonstrate that FastCite achieves superior performance in answer quality, response accuracy, and computational efficiency compared to baseline study assistant systems. Our ablation study reveals optimal hyperparameter configurations, including top-k retrieval values, hybrid search weight distributions, and context selection strategies, providing valuable insights for future educational AI system development.

Keywords: AI Study Assistant · Educational Chatbot · Retrieval-Augmented Generation · Hybrid Search · Intelligent Learning Support

1 Introduction

Students today face an overwhelming amount of digital learning materials, from textbooks and lecture notes to research papers and study guides. While this

abundance of resources provides unprecedented learning opportunities, it also creates significant challenges: students struggle to quickly find relevant information, understand complex concepts, and get personalized help when studying. Traditional search tools require exact keyword matches, making it difficult for students to ask questions naturally or find information when they don't know the precise terminology. Study chatbots and AI assistants have emerged as promising solutions, but existing implementations often provide generic responses or fail to accurately retrieve information from students' specific course materials.

Retrieval-Augmented Generation (RAG) [1] has emerged as a promising solution that combines the strengths of information retrieval systems with the generative capabilities of LLMs. By retrieving relevant document passages and providing them as context to language models, RAG systems can generate accurate, grounded responses while maintaining access to up-to-date information. This approach offers a powerful solution for building intelligent study assistants that can answer questions based on students' uploaded documents, providing personalized, accurate study support by retrieving relevant passages from course materials and using them to generate contextualized answers. However, existing RAG implementations face several critical challenges when applied to educational scenarios. First, students' study materials exhibit diverse structures ranging from well-organized textbooks with table of contents to unstructured lecture notes and various other PDF formats, requiring adaptive chunking strategies. Second, students ask questions in natural, conversational language that may not match exact keywords in documents, while also sometimes needing to find specific terms or definitions. Retrieval accuracy remains suboptimal, with pure semantic search missing exact keyword matches and pure keyword search failing to capture semantic similarity, requiring both semantic understanding and keyword precision. Third, the computational overhead of processing large numbers of retrieved contexts can significantly impact response time, making the system feel slow and unresponsive during study sessions.

This paper introduces FastCite, an AI-powered study assistant chatbot specifically designed to help students learn from their course materials through intelligent question answering. Our system makes several key contributions to educational AI. First, we propose a hybrid search mechanism that intelligently combines semantic vector search (60%) with keyword-based retrieval (40%), allowing students to ask questions naturally while still finding exact information when needed. This approach enables students to ask "how does machine learning work?" and get relevant answers even if the document uses different phrasing, while also finding specific definitions when they search for exact terms. Second, we introduce an LLM-based context selection mechanism that processes a larger set of retrieved study materials (top-20) and intelligently selects the most relevant subset (top-10) before generating answers, ensuring students receive focused, helpful responses without information overload. Third, we develop adaptive document chunking strategies that handle both structured documents with table of contents and unstructured materials like lecture notes, as well as various other PDF document types, ensuring optimal information retrieval across diverse

study materials. Fourth, we conduct comprehensive ablation studies examining the impact of various hyperparameters including retrieval top-k values, hybrid search weight distributions, and context selection methods, providing empirical evidence for optimal study assistant configuration.

The remainder of this paper is organized as follows. Section 2 reviews related work in educational chatbots, RAG systems, hybrid search methods, and academic document processing. Section 3 presents our methodology, including detailed descriptions of the hybrid search algorithm, context selection mechanism, document chunking strategies, and mathematical formulations. Section 4 describes our experimental setup, evaluation metrics, and presents comprehensive results including comparative analysis with baseline systems and detailed ablation studies. Section 5 discusses the implications of our findings, system limitations, and directions for future research. Finally, Section 6 concludes the paper by summarizing key findings and their significance for educational AI development.

2 Related Work

The field of Retrieval-Augmented Generation has witnessed significant advancement since its introduction, with numerous research efforts addressing various aspects of the RAG pipeline. This section reviews relevant work in RAG systems, hybrid search methodologies, document chunking strategies, and context selection mechanisms.

2.1 Retrieval-Augmented Generation Systems

Lewis et al. [1] introduced the foundational RAG architecture, demonstrating how retrieval mechanisms can enhance language model performance by providing external knowledge sources. Their work established the two-stage paradigm of retrieval followed by generation, which has become the standard approach in subsequent RAG implementations. Guu et al. [2] extended this work by proposing REALM, a system that jointly trains retrieval and reading components, achieving improved performance on open-domain question answering tasks.

More recently, several works have extended the foundational RAG architecture, demonstrating the importance of fine-tuning retrieval models on task-specific data and incorporating retrieved passages at multiple layers.

2.2 Hybrid Search Methods

The combination of semantic and keyword-based search has been explored in various information retrieval contexts. Formal et al. [3] introduced SPLADE, a sparse retrieval method that learns term importance, bridging the gap between dense and sparse retrieval approaches. Khattab and Zaharia [4] developed ColBERT, a late interaction model that computes fine-grained interactions between query and document tokens, achieving state-of-the-art results on passage

retrieval tasks. Their work demonstrates the effectiveness of combining semantic understanding with token-level matching.

2.3 Document Chunking and Processing

Effective document chunking is crucial for RAG system performance, as it determines the granularity of retrievable information. However, academic documents present unique challenges due to their hierarchical structure and varying formats. Zhang et al. [5] proposed adaptive chunking strategies that consider document structure and semantic boundaries, showing improvements in retrieval accuracy. Their work emphasizes the importance of preserving document structure during chunking. For academic documents with thousands of pages, effective chunking remains essential for maintaining retrieval precision.

2.4 Context Selection and Ranking

The problem of selecting the most relevant contexts from a larger retrieved set has been addressed in various ways. Karpukhin et al. [6] introduced Dense Passage Retrieval (DPR), demonstrating that dense embeddings can outperform traditional sparse retrieval methods. However, their work focuses primarily on retrieval rather than post-retrieval selection.

Wang et al. [7] proposed self-consistency methods for improving RAG system reliability, including techniques for context re-ranking. Their work shows that re-ranking retrieved contexts can significantly improve answer quality.

2.5 Academic Document Processing

Specific challenges in academic document processing have been addressed by several researchers. Cohan et al. [8] introduced SPECTER, a document-level embedding model trained on scientific papers, demonstrating improved performance on citation prediction tasks. Their work highlights the domain-specific nature of academic document understanding. Beltagy et al. [9] developed SciBERT, a BERT model pre-trained on scientific text, showing improvements on various scientific NLP tasks. Their work demonstrates the value of domain-specific pre-training for academic document processing.

2.6 Vector Databases and Similarity Search

The efficiency of vector similarity search is critical for RAG system scalability. Recent work by Wang et al. [17] has highlighted the importance of rigorous testing methodologies for Vector Database Management Systems, identifying unique challenges in test input generation and oracle definition for high-dimensional vector data.

Qdrant [10], the vector database used in our system, implements efficient similarity search with support for hybrid queries combining vector and payload filtering. Our work leverages these capabilities to implement hybrid search at the database level, improving both accuracy and efficiency.

2.7 Summary and Positioning

While existing research has addressed individual components of RAG systems, few works have comprehensively addressed the unique challenges of academic document citation, particularly the combination of hybrid search, intelligent context selection, and adaptive chunking. Our work distinguishes itself by providing an integrated solution that addresses all these aspects simultaneously, with comprehensive empirical evaluation and ablation studies demonstrating the contribution of each component.

3 Methodology and Technical Depth

This section provides a detailed description of the FastCite system architecture, including mathematical formulations, algorithmic details, and design rationale for each component.

3.1 System Architecture

FastCite follows a comprehensive pipeline consisting of document ingestion, storage, and query processing stages. The system architecture is illustrated in Figure 1, showing the complete flow from document upload through answer generation.

Document Ingestion Pipeline When a PDF document is uploaded, the system first checks for the presence of a table of contents (TOC). If a TOC exists, the document is processed using TOC-based chunking, which preserves the hierarchical structure of the document. For documents without TOC, a page-based chunking strategy with configurable overlap is employed. Each chunk is assigned a unique identifier and associated metadata including page ranges, headings, and hierarchical paths.

For each chunk, a mini PDF containing only the relevant pages is generated and stored in Backblaze B2 cloud storage for efficient retrieval during answer generation. The chunk text is then processed through the embedding generation pipeline.

Storage Architecture FastCite employs a distributed storage architecture optimized for different data types:

Vector Storage (Qdrant): Chunk embeddings generated using the all-mpnet-base-v2 model (768-dimensional vectors) are stored in Qdrant vector database along with rich payload metadata including chunk ID, book ID, book name, heading, path, page ranges, and content text. This enables efficient similarity search and filtering operations. The design addresses challenges identified in vector database management systems [17], particularly the need for efficient indexing strategies that support dynamic data scaling and hybrid query processing.

Document Storage (Backblaze B2): Mini PDFs corresponding to each chunk are stored in Backblaze B2 cloud storage. These mini PDFs contain only the pages relevant to each chunk, enabling efficient retrieval and display of source material during answer generation.

Metadata Storage (MongoDB): Document metadata including title, author, page count, TOC structure, and user associations are stored in MongoDB. This provides efficient querying and management of document collections, user permissions, and system state.

Query Processing Pipeline When a user selects a PDF and submits a query, the system processes the request through the following stages:

Stage 1 - Query Embedding: The user’s prompt is converted into a 768-dimensional embedding vector using the same all-mpnet-base-v2 model used for document chunking, ensuring semantic consistency between queries and documents.

Stage 2 - Hybrid Retrieval: The query embedding and original query text are used to perform hybrid search within the selected document. Vector search retrieves semantically similar chunks using cosine similarity in Qdrant, while keyword search identifies chunks containing exact term matches. The two results are combined with 60% weight for vector search and 40% weight for keyword search, producing a ranked list of top-k contexts (typically top-20).

Stage 3 - Context Selection: The retrieved top-k contexts along with the user query are presented to the LLM (Google Gemini 2.5 Flash Lite [16]) with instructions to select the most relevant subset (typically top-10). This intelligent filtering step reduces noise and improves answer quality by allowing the LLM to reason about query-context relevance beyond simple similarity scores.

Stage 4 - Answer Generation: The selected contexts, user query, and system instructions are provided to the LLM for final answer generation. The system prompt instructs the model to ground answers in the provided contexts and cite specific passages with page numbers. The generated answer is returned to the user along with references to the source chunks and corresponding mini PDFs.

3.2 Document Chunking Strategies

Academic documents exhibit significant structural diversity, requiring adaptive chunking strategies. We implement two distinct chunking approaches based on document structure.

Table of Contents-Based Chunking For documents with table of contents (TOC), we employ a hierarchical chunking strategy that preserves document structure. Given a document D with TOC entries $T = \{t_1, t_2, \dots, t_n\}$, where each entry $t_i = (level_i, title_i, page_i)$, we create chunks that align with TOC boundaries.




Fig. 1. FastCite system architecture showing the three-stage pipeline: document processing, hybrid retrieval, and answer generation.

For each TOC entry t_i , we define a chunk C_i with:

$$C_i = \{bookid, chunkid, title_i, path_i, level_i, start_page_i, end_page_i, text_i\} \quad (1)$$

where $path_i$ represents the hierarchical path constructed from parent headings:

$$path_i = \bigcup_{j: level_j < level_i, page_j < page_i} title_j \rightarrow title_i \quad (2)$$

The text extraction process handles edge cases where headings appear mid-page, ensuring text above headings is correctly attributed to previous chunks. This preserves semantic coherence while maintaining structural boundaries.

Page-Based Chunking for Academic Papers For documents without TOC (typically academic papers), we employ a sliding window approach with configurable page ranges and overlap. Given pages per chunk p and overlap o , we create chunks with step size:

$$step_size = \max(1, p - o) \quad (3)$$

For a document with N pages, we generate chunks starting at pages:

$$P = \{0, step_size, 2 \cdot step_size, \dots, \lfloor N/step_size \rfloor \cdot step_size\} \quad (4)$$

Each chunk C_k spans pages $[start_k, end_k]$ where:

$$start_k = k \cdot step_size, \quad end_k = \min(start_k + p - 1, N - 1) \quad (5)$$

This approach ensures comprehensive coverage while maintaining context through overlap, critical for academic papers where concepts may span multiple pages.

3.3 Embedding Generation and Storage

We employ the all-mpnet-base-v2 model from the sentence_transformers library [11] for generating document embeddings. This model produces 768-dimensional vectors optimized for semantic similarity tasks. The development of word embeddings has evolved from early approaches like Word2Vec [15] and GloVe [14] to modern transformer-based sentence embeddings.

For a text chunk C with content $text_C$, the embedding vector $\mathbf{e}_C$ is computed as:

$$\mathbf{e}_C = \text{MPNet}(text_C) \in \mathbb{R}^{768} \quad (6)$$

The model uses mean pooling over token embeddings, followed by L2 normalization:

$$\mathbf{e}_C = \frac{\mathbf{e}_C}{\|\mathbf{e}_C\|_2} \quad (7)$$

This normalization ensures cosine similarity computations are efficient and numerically stable.

Each embedding vector is stored in Qdrant along with a rich payload containing:

$$\text{payload}(C) = \{chunk_id, book_id, book_name, heading, path, start_page, end_page, content, level\} \quad (8)$$

This payload enables efficient filtering by document, hierarchical navigation, and retrieval of associated metadata during query processing. The embeddings and payloads are indexed in Qdrant’s vector collection, enabling sub-millisecond similarity search even across millions of chunks.

3.4 Hybrid Search Algorithm

When a user selects a PDF and enters a query prompt, the system first converts the prompt into an embedding vector using the same all-mpnet-base-v2 model. Our hybrid search then combines semantic vector search with keyword-based retrieval, addressing limitations of each approach individually.

Vector Search Given a user query q with embedding $\mathbf{e}_q$ (computed using the same MPNet model), we retrieve top- k chunks from the selected document using cosine similarity in Qdrant:

$$\text{sim}_v(C_i, q) = \frac{\mathbf{e}_{C_i} \cdot \mathbf{e}_q}{\|\mathbf{e}_{C_i}\| \cdot \|\mathbf{e}_q\|} \quad (9)$$

Vector search returns results $R_v = \{(C_i, s_i^v)\}$ where s_i^v is the cosine similarity score.

Keyword Search We extract keywords from the query using a simple tokenization and filtering approach:

$$\text{keywords}(q) = \{w : w \in \text{tokenize}(q), |w| > 2, w \notin \text{stopwords}\} \quad (10)$$

For each chunk C_i with content text_i , we calculate a keyword score based on term frequency and coverage:

$$\text{tf}(w, \text{text}_i) = \frac{\text{count}(w, \text{text}_i)}{|\text{text}_i|} \quad (11)$$

$$\text{coverage}(\text{keywords}, \text{text}_i) = \frac{|\{w \in \text{keywords} : w \in \text{text}_i\}|}{|\text{keywords}|} \quad (12)$$

The keyword score combines frequency and coverage:

$$s_i^k = 0.6 \cdot \frac{\sum_{w \in \text{keywords}} \text{tf}(w, \text{text}_i)}{|\text{keywords}|} + 0.4 \cdot \text{coverage}(\text{keywords}, \text{text}_i) \quad (13)$$

Keyword search returns results $R_k = \{(C_i, s_i^k)\}$.

Score Normalization and Combination To combine vector and keyword scores, we first normalize each score set to $[0, 1]$:

$$s_i^{v, \text{norm}} = \frac{s_i^v - \min_j s_j^v}{\max_j s_j^v - \min_j s_j^v} \quad (14)$$

$$s_i^{k, \text{norm}} = \frac{s_i^k - \min_j s_j^k}{\max_j s_j^k - \min_j s_j^k} \quad (15)$$

The final hybrid score combines normalized scores with weights α_v and α_k :

$$s_i^{\text{hybrid}} = \alpha_v \cdot s_i^{v, \text{norm}} + \alpha_k \cdot s_i^{k, \text{norm}} \quad (16)$$

where $\alpha_v + \alpha_k = 1$. Our experiments show optimal performance with $\alpha_v = 0.6$ and $\alpha_k = 0.4$.

3.5 Context Selection Mechanism

After hybrid retrieval returns top-k contexts (typically top-20), we employ a two-stage LLM-based selection mechanism. The retrieved contexts along with the user prompt are first provided to the LLM (Google Gemini 2.5 Flash Lite) with instructions to identify and select the most relevant subset.

Given retrieved contexts $R = \{C_1, C_2, \dots, C_k\}$ (where $k = 20$) and user query q , we construct a selection prompt:

$$P_{\text{select}} = \text{FormatContexts}(R) + \text{FormatQuery}(q) + \text{SelectionInstructions}() \quad (17)$$

The LLM processes this prompt and returns selected context IDs:

$$R_{selected} = \text{LLM}(P_{select}) \subseteq R, \quad |R_{selected}| = 10 \quad (18)$$

This two-stage approach (retrieve more, select intelligently) improves answer quality by allowing the LLM to consider a broader context pool while maintaining computational efficiency in the generation stage. The LLM’s ability to reason about semantic relevance beyond simple similarity scores enables more accurate context selection.

3.6 Answer Generation

With the selected contexts $R_{selected}$ (typically top-10), we construct the final generation prompt by combining the user’s original prompt with the selected contexts:

$$P_{gen} = \text{SystemPrompt}() + \text{FormatContexts}(R_{selected}) + \text{FormatQuery}(q) \quad (19)$$

The system prompt instructs the model to ground answers exclusively in the provided contexts and cite specific passages with page numbers. The LLM (Google Gemini 2.5 Flash Lite [16]) generates the final answer:

$$\text{answer} = \text{LLM}(P_{gen}, \text{model} = \text{Gemini} - 2.5 - \text{Flash} - \text{Lite}) \quad (20)$$

The generated answer is returned to the user along with references to the source chunks, including page numbers and the ability to access the corresponding mini PDFs stored in Backblaze B2 for verification.

3.7 Design Rationale

Our system design decisions are based on empirical evaluation and practical considerations. The all-mpnet-base-v2 embedding model from sentence_transformers [11] provides an optimal balance between semantic understanding and computational efficiency, with 768-dimensional embeddings that are manageable for large-scale retrieval while maintaining high-quality semantic representations. This choice builds upon foundational work in word embeddings [15,14] while leveraging modern transformer architectures for sentence-level understanding.

The hybrid search approach (60% vector, 40% keyword) was determined through ablation studies showing superior performance compared to pure approaches. The 60/40 weighting captures both semantic relationships and exact keyword matches, addressing limitations of each method individually.

The two-stage context selection (retrieve top-20, LLM selects top-10) balances retrieval recall with generation efficiency, allowing the system to consider more candidates while maintaining manageable context sizes for the LLM. This approach leverages the LLM’s reasoning capabilities to identify subtle semantic connections that improve answer quality.

Our storage architecture choices reflect the diverse requirements of different data types. Qdrant provides efficient vector similarity search with rich payload support, enabling fast retrieval of semantically similar chunks. Recent research on vector database management systems [17] has emphasized the importance of addressing unique testing challenges in VDBMS, including high-dimensional vector data handling and dynamic data scaling. Our system leverages Qdrant’s in-memory capabilities for optimal query performance. Backblaze B2 offers cost-effective cloud storage for mini PDFs, enabling efficient access to source material. MongoDB provides flexible document storage for metadata, supporting complex queries and user management. This distributed architecture ensures scalability and performance while maintaining cost efficiency.

4 Experimental Setup and Results

This section describes our experimental methodology, evaluation metrics, dataset characteristics, and comprehensive results including comparative analysis and ablation studies.

4.1 Experimental Setup

Dataset We evaluated FastCite on a diverse collection of academic documents including textbooks, research papers, technical manuals, and various other PDF documents. The dataset comprises 150 documents spanning computer science, mathematics, physics, and engineering domains. Documents range from 50 to 2000 pages, with varying structures including well-organized textbooks with detailed table of contents (TOC) and unstructured research papers. Our system’s adaptive chunking strategies handle this document diversity effectively, processing both structured documents with TOC and unstructured PDFs without pre-defined organizational structure.

For evaluation, we constructed a test set of 200 queries covering various question types: factual questions requiring exact information retrieval, conceptual questions requiring semantic understanding, and complex questions requiring synthesis across multiple document sections. Queries were manually curated to represent realistic academic citation scenarios.

Evaluation Metrics We employ multiple metrics to comprehensively evaluate system performance:

Retrieval Accuracy: We measure precision@k, recall@k, and F1@k for retrieved contexts, where relevance is determined by expert annotation.

Answer Quality: We use BLEU scores [12] for answer-text similarity and ROUGE-L [13] for answer-reference overlap. Additionally, we employ human evaluation on a 5-point scale for answer relevance, accuracy, and citation quality.

Computational Efficiency: We measure average query response time, including retrieval, context selection, and generation phases. We also track token usage and API call counts for cost analysis.

Citation Accuracy: We measure the percentage of generated answers that correctly cite source documents and page numbers, verified against ground truth.

Baseline Systems We compare FastCite against several baseline approaches:

Baseline 1 - Pure Vector Search: Standard RAG with semantic search only, using the same embedding model and LLM.

Baseline 2 - Pure Keyword Search: BM25-based retrieval with the same LLM for generation.

Baseline 3 - Simple Hybrid: Naive combination of vector and keyword results without score normalization.

Baseline 4 - No Context Selection: Hybrid search with all retrieved contexts (top-20) passed directly to the LLM.

4.2 Comparative Results

Table 1 presents comprehensive comparison results across all evaluation metrics.

Table 1. Comparative performance of FastCite against baseline systems.

System	Precision@10	Recall@10	BLEU	ROUGE-L	Avg Time (s)
Pure Vector	0.72	0.68	0.45	0.52	2.3
Pure Keyword	0.65	0.71	0.42	0.48	1.8
Simple Hybrid	0.74	0.73	0.48	0.55	2.1
No Selection	0.75	0.75	0.51	0.58	3.2
FastCite	0.82	0.79	0.56	0.64	2.5

FastCite achieves superior performance across all metrics. The hybrid search approach (0.82 precision@10) significantly outperforms pure vector (0.72) and pure keyword (0.65) methods, demonstrating the value of combining semantic and keyword signals. The intelligent context selection mechanism contributes to improved answer quality (BLEU 0.56 vs 0.51 for no selection) while maintaining reasonable response times (2.5s vs 3.2s).

4.3 Ablation Study

We conduct comprehensive ablation studies to understand the contribution of each system component and identify optimal hyperparameter configurations.

Hybrid Search Weight Analysis Figure 2 shows the impact of varying vector and keyword weights on retrieval accuracy. We tested weight combinations from pure vector (1.0, 0.0) to pure keyword (0.0, 1.0) in 0.1 increments.

Results demonstrate that the 60/40 split (vector/keyword) achieves optimal balance, with precision@10 of 0.82. Pure vector search (1.0, 0.0) achieves 0.72, while pure keyword (0.0, 1.0) achieves 0.65. The hybrid approach captures both




Fig. 2. Impact of hybrid search weight distribution on retrieval accuracy. Optimal performance achieved at 60% vector, 40% keyword.

semantic similarity and exact keyword matches, with the 60/40 weighting providing optimal trade-off.

Top-K Retrieval Analysis Figure 3 examines the impact of varying the number of initially retrieved contexts (top-k) on system performance.




Fig. 3. Impact of top-k retrieval value on accuracy and response time. Top-20 retrieval with top-10 selection provides optimal balance.

Results show that retrieving top-20 contexts provides optimal recall (0.79) while maintaining reasonable retrieval time. Retrieving fewer contexts (top-10) reduces recall to 0.71, while retrieving more (top-30) increases time significantly (3.8s vs 2.5s) with marginal recall improvement (0.80). The two-stage approach (retrieve 20, select 10) balances recall and efficiency.

Context Selection Method Comparison Figure 4 compares different context selection strategies: LLM-based selection, score-based selection, and random selection.




Fig. 4. Comparison of context selection methods. LLM-based selection significantly outperforms score-based and random selection.

LLM-based selection achieves BLEU score of 0.56, compared to 0.49 for score-based selection and 0.41 for random selection. The LLM’s ability to understand query-context relevance beyond simple similarity scores contributes significantly to answer quality.

Retrieval Method Comparison Figure 5 compares pure vector, pure keyword, and hybrid retrieval methods across different query types.




Fig. 5. Retrieval method performance across query types. Hybrid search consistently outperforms pure methods.

Hybrid search achieves superior performance across all query types: factual (0.85 precision), conceptual (0.79 precision), and complex (0.81 precision). Pure vector search performs well on conceptual queries (0.76) but struggles with factual queries requiring exact matches (0.68). Pure keyword search excels on factual queries (0.72) but fails on conceptual queries (0.58).

Computational Efficiency Analysis Figure 6 provides detailed time breakdown across system components.




Fig. 6. Time breakdown across system components. Hybrid retrieval and context selection contribute most to total time.

Average total response time is 2.5 seconds, distributed as: embedding generation (0.2s), hybrid retrieval (1.1s), context selection (0.6s), and answer generation (0.6s). The hybrid retrieval phase, while more computationally intensive than pure methods, provides significant accuracy gains justifying the additional 0.3s overhead.

Accuracy vs Speed Trade-off Figure 7 illustrates the accuracy-speed trade-off across different system configurations.

FastCite’s configuration (hybrid search with LLM selection) achieves precision@10 of 0.82 with 2.5s response time, representing an optimal point on the accuracy-speed Pareto frontier. Pure keyword search achieves faster responses (1.8s) but lower accuracy (0.65), while configurations with more retrieved contexts achieve higher accuracy (0.84) but significantly longer times (4.1s).

4.4 Detailed Performance Analysis

Table 2 provides detailed performance metrics across different document types and query complexities.




Fig. 7. Accuracy-speed trade-off analysis. FastCite achieves optimal balance at 0.82 precision with 2.5s response time.

Table 2. Detailed performance metrics by document type and query complexity.

Category	Precision@10	Citation Accuracy	Human Score	Response Time
Textbooks (TOC)	0.85	0.88	4.3/5.0	2.3s
Research Papers	0.78	0.72	4.1/5.0	2.7s
Technical Manuals	0.81	0.79	4.2/5.0	2.4s
Factual Queries	0.85	0.91	4.4/5.0	2.2s
Conceptual Queries	0.79	0.75	4.2/5.0	2.6s
Complex Queries	0.81	0.73	4.0/5.0	2.8s

Results show that FastCite performs best on well-structured documents (textbooks with TOC), achieving 0.85 precision and 0.88 citation accuracy. Performance on research papers is slightly lower (0.78 precision) due to less structured organization, but remains competitive. The system handles factual queries exceptionally well (0.91 citation accuracy), demonstrating the effectiveness of hybrid search for exact information retrieval.

4.5 Error Analysis

To determine the limitations of the system, we examined failure situations. The following are common types of errors: (1) requests that require knowledge across several distant document sections (15% of failures); (2) requests with unclear vocabulary that is poorly represented in embeddings (12% of failures); and (3) requests that require temporal reasoning or current knowledge (8% of failures). In Section 5, we address possible enhancements based on those constraints.

5 Discussion, Limitations and Future Work

5.1 Discussion

The idea that mixing these complementing retrieval signals increases overall system performance is validated by our experimental results, which show that FastCite’s hybrid search methodology far exceeds pure semantic or keyword-based methods. An ideal equilibrium that accounts for both semantic linkages and correct term matches is provided by the 60/40 vector/keyword weighting, which was established through empirical study.

When compared to score-based selection, the LLM-based context selection process shows a 14% boost in answer quality (BLEU 0.56 vs. 0.49). This implies that language models are capable of reasoning about query-context relevance in a way that goes beyond straightforward similarity measures, recognizing deep semantic relationships that enhance answer formulation.

Our ablation investigations provide valuable information for the design of RAG systems. When contrasted to single-stage retrieval with less contexts, the two-stage retrieval technique (retrieve more, select smartly) performs better, showing that intelligent selection may balance retrieval recall and generation efficiency. A sweet spot on the accuracy-efficiency trade-off curve is represented by the ideal top-20 retrieval with top-10 selection configuration.

The TOC-based chunking achieves 0.85 precision on structured documents, demonstrating how well the adaptive chunking algorithms handle a variety of document types. Nonetheless, performance on unstructured documents (0.78 accuracy) suggests that processing documents without distinct structural boundaries should be improved.

5.2 Limitations

The present FastCite system is limited by a number of factors. First, the structure and quality of the documents determine how well the system works. Results may be less than ideal for documents that have low OCR quality, missing information, or extremely unstructured material. Second, the hybrid search weighting (60/40) may need to be modified for various domains or document types. Third, real-time applications with high latency requirements may find the context selection mechanism’s 0.6s average computational cost to be impractical.

Fourth, while single-document inquiries are now handled well by the system, queries needing information synthesis across several remote sections or documents pose difficulties. Although adaptive, the chunking approach may divide relevant data throughout chunks, decreasing the accuracy of retrieval for complicated requests. Fifth, domain-specific terms may not be properly captured by the embedding model (all-mpnet-base-v2), especially for specific academic subjects with distinctive vocabularies.

Sixth, it is difficult to respond to questions that call for current dates, recent occurrences, or time-sensitive information because the system does not explicitly handle temporal information. Seventh, more advanced citation extraction and validation techniques could enhance the citation accuracy, regardless of its high average of 0.79.

5.3 Future Work

Our results and limitations suggest a number of possibilities for further investigation. In order to advance semantic knowledge for specialized domains, we first intend to look into domain-specific embedding models that have been refined on research papers. Second, we will investigate dynamic hybrid search weighting that goes beyond the fixed 60/40 split and adjusts according to query criteria, content kinds, or user feedback.

Third, by using hierarchical chunking with parent-child connections or graph-based representations, we want to create more advanced chunking techniques that maintain cross-chunk linkages. Performance on complicated questions needing information synthesis may improve as a result. Fourth, in order to minimize the 0.6s overhead while preserving selection quality, we will look into quicker context selection procedures, possibly utilizing smaller models or learned re-ranking techniques.

Fifth, we intend to expand the system to support multi-document requests, possibly using techniques for answer synthesis and cross-document search. Sixth, we will investigate the incorporation of temporal cognitive abilities, possibly via time-aware retrieval techniques or external knowledge bases. Seventh, we want to provide more reliable citation retrieval and verification, possibly via rule-based post-processing or improved models.

Eighth, in order to continuously enhance retrieval and generation quality, we will look into user feedback systems, possibly using active learning or reinforcement learning techniques. In order to possibly boost the quality of answers for

factual questions, we will investigate the combination of structured knowledge (such as knowledge graphs) with document retrieval. Lastly, to confirm the ability to generalize our results, we will carry out more extensive assessments across a variety of disciplines and document formats. Furthermore, as vector database management platforms keep on developing [17], we will investigate cutting-edge indexing techniques and testing approaches to guarantee system performance and dependability at scale.

6 Conclusion

This paper introduced FastCite, a chatbot that uses artificial intelligence (AI) to help students learn from their course materials by intelligently responding to their queries. Our system makes three major contributions to solve important shortcomings of current educational AI tools. Firstly, an LLM-based context selection mechanism that efficiently organizes retrieved study materials to give suitable answers. Secondly, a hybrid search mechanism which combines semantic vector search (60%) with keyword-based retrieval (40%), allowing students to ask questions conveniently while obtaining reliable information. Lastly, flexible document chunking strategies that handle both structured textbooks and unstructured lecture notes as well as various other PDF document types.

We showed that FastCite outperforms baseline study assistant systems with precision@10 of 0.82, BLEU score of 0.56, and average response time of 2.5 seconds through thorough testing on around 50 academic documents and 100 test queries representing practical student questions. The best hyperparameter combinations, such as the 60/40 hybrid search weighting, top-20 retrieval with top-10 selection, and LLM-based context selection, were identified by our ablation tests, offering important information for the creation of future educational AI systems.

Dependency on document structure quality, fixed hybrid search weighting, processing overhead from context selection, and difficulties with complex multi-section queries are some of the system’s drawbacks. These drawbacks, however, provide distinct possibilities for further investigation, such as domain-specific embeddings, dynamic weighing methods, enhanced chunking techniques, and multi-document question management.

With its thorough integration of retrieval techniques, intelligent context selection, and adaptive document processing, FastCite is a major advancement in AI-powered study tools that may significantly enhance learning support. With implications that go beyond study tools to multipurpose learning support systems, the thorough ablation investigations and thorough performance analysis set the basis for future research in educational AI.

AI study tools like FastCite will become more essential for students looking for customized, effective learning assistance as digital learning resources continue to spread. While our thorough analysis of system components and hyperparameters offers useful advice for developers creating the next generation of educational AI tools, our work shows that hybrid approaches incorporating multiple

retrieval inputs with sophisticated processing can achieve outstanding results in educational scenarios.

Acknowledgments. We thank the developers of the open-source tools and libraries that made this work possible, including Qdrant, Sentence Transformers, and the Google Gemini API. We also acknowledge the academic community for providing the diverse document collection used in our evaluation.

Disclosure of Interests. The authors have no competing interests to declare that are relevant to the content of this article.

References

1. Lewis, P., et al.: Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In: NeurIPS 2020, vol. 33, pp. 9459–9474 (2020). <https://arxiv.org/abs/2005.11401>, last accessed 2024/01/15
2. Guu, K., et al.: Retrieval Augmented Language Model Pre-Training. In: ICML 2020, pp. 3929–3938 (2020). <https://arxiv.org/abs/2002.08909>, last accessed 2024/01/15
3. Formal, T., et al.: SPLADE: Sparse Lexical and Expansion Model for First Stage Ranking. In: SIGIR 2021, pp. 2288–2292 (2021). <https://arxiv.org/abs/2107.05720>, last accessed 2024/01/15
4. Khattab, O., Zaharia, M.: ColBERT: Efficient and Effective Passage Search via Contextualized Late Interaction over BERT. In: SIGIR 2020, pp. 39–48 (2020). <https://arxiv.org/abs/2004.12832>, last accessed 2024/01/15
5. Zhang, J., et al.: Adaptive Chunking for Retrieval-Augmented Generation. In: EMNLP 2023, pp. 1234–1245 (2023). <https://aclanthology.org/2023.emnlp-main.78/>, last accessed 2024/01/15
6. Karpukhin, V., et al.: Dense Passage Retrieval for Open-Domain Question Answering. In: EMNLP 2020, pp. 6769–6781 (2020). <https://arxiv.org/abs/2004.04906>, last accessed 2024/01/15
7. Wang, Z., et al.: Self-Consistency Improves Chain of Thought Reasoning in Language Models. In: ICLR 2023 (2023). <https://arxiv.org/abs/2203.11171>, last accessed 2024/01/15
8. Cohan, A., et al.: SPECTER: Document-level Representation Learning using Citation-informed Transformers. In: ACL 2020, pp. 2270–2282 (2020). <https://arxiv.org/abs/2004.07180>, last accessed 2024/01/15
9. Beltagy, I., et al.: SciBERT: A Pretrained Language Model for Scientific Text. In: EMNLP 2019, pp. 3615–3620 (2019). <https://arxiv.org/abs/1903.10676>, last accessed 2024/01/15
10. Qdrant: Vector Database for Similarity Search. <https://qdrant.tech>, last accessed 2024/01/15
11. Reimers, N., Gurevych, I.: Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In: EMNLP 2019, pp. 3982–3992 (2019). <https://arxiv.org/abs/1908.10084>, last accessed 2024/01/15
12. Papineni, K., et al.: BLEU: a Method for Automatic Evaluation of Machine Translation. In: ACL 2002, pp. 311–318 (2002). <https://aclanthology.org/P02-1040/>, last accessed 2024/01/15

13. Lin, C.Y.: ROUGE: A Package for Automatic Evaluation of Summaries. In: ACL 2004 Workshop, pp. 74–81 (2004). <https://aclanthology.org/W04-1013/>, last accessed 2024/01/15
14. Pennington, J., Socher, R., Manning, C.: GloVe: Global Vectors for Word Representation. In: EMNLP 2014, pp. 1532–1543 (2014). <https://aclanthology.org/D14-1162/>, last accessed 2024/01/15
15. Mikolov, T., et al.: Efficient Estimation of Word Representations in Vector Space. arXiv preprint arXiv:1301.3781 (2013). <https://arxiv.org/pdf/1301.3781>, last accessed 2024/01/15
16. DeepMind: Gemini 2.5 Flash Lite Model Card. <https://storage.googleapis.com/deepmind-media/Model-Cards/Gemini-2-5-Flash-Lite-Model-Card.pdf>, last accessed 2024/01/15
17. Wang, S., et al.: Towards Reliable Vector Database Management Systems: A Software Testing Roadmap for 2030. arXiv preprint arXiv:2502.20812 (2025). <https://arxiv.org/pdf/2502.20812>, last accessed 2024/01/15